

mk 3/08/2023

```
In [494]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [495]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2004.csv")
df
```

Out[495]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10
0	2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.990
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950
2	2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.480
3	2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.070
4	2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.080
...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.689
245493	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.840
245494	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.630
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.389

245496 rows × 17 columns



```
In [496]: df=df.dropna()
```

```
In [556]: df=df.head(1000)
```

```
In [557]: df.columns
```

```
Out[557]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
               dtype='object')
```

```
In [558]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 20 entries, 5 to 184  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        20 non-null    object  
1   BEN         20 non-null    float64  
2   CO          20 non-null    float64  
3   EBE         20 non-null    float64  
4   MXY         20 non-null    float64  
5   NMHC        20 non-null    float64  
6   NO_2        20 non-null    float64  
7   NOx         20 non-null    float64  
8   OXY         20 non-null    float64  
9   O_3         20 non-null    float64  
10  PM10        20 non-null    float64  
11  PM25        20 non-null    float64  
12  PXY         20 non-null    float64  
13  SO_2        20 non-null    float64  
14  TCH         20 non-null    float64  
15  TOL         20 non-null    float64  
16  station     20 non-null    int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 2.8+ KB
```


```
In [559]: data=df[['BEN', 'TOL', 'PXY']]  
data
```

Out[559]:

	BEN	TOL	PXY
5	3.24	16.93	4.16
22	0.55	2.53	0.39
26	1.80	7.92	2.21
32	1.94	8.18	2.14
49	0.29	1.52	0.34
53	1.23	5.00	1.58
59	1.64	6.35	1.63
76	0.29	1.23	0.28
80	1.11	3.98	1.32
86	2.45	9.26	2.30
103	0.27	1.01	0.37
107	1.20	4.03	1.47
113	2.78	9.86	2.50
130	0.29	0.91	0.87
134	1.19	3.97	1.53
140	1.04	3.84	1.02
157	0.26	0.88	0.74
161	1.01	3.19	1.16
167	1.34	4.92	1.17
184	0.29	0.86	1.00


```
In [560]: data.plot.line(subplots=True)
```

Out[560]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)



```
In [561]: data.plot.line()
```


```
Out[561]: <AxesSubplot:>
```



```
In [562]: b=data[0:50]
```


```
In [563]: b.plot.bar()
```

```
Out[563]: <AxesSubplot:>
```




```
In [564]: data.plot.hist()
```

```
Out[564]: <AxesSubplot:ylabel='Frequency'>
```




```
In [565]: data.plot.area()
```

```
Out[565]: <AxesSubplot:>
```




In [566]: `data.plot.box()`

Out[566]: `<AxesSubplot:>`




In [567]: `b.plot.pie(y='BEN')`

Out[567]: `<AxesSubplot:ylabel='BEN'>`



```
In [568]: data.plot.scatter(x='TOL', y='BEN')
```

```
Out[568]: <AxesSubplot:xlabel='TOL', ylabel='BEN'>
```




```
In [569]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20 entries, 5 to 184
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        20 non-null    object
1   BEN         20 non-null    float64
2   CO          20 non-null    float64
3   EBE         20 non-null    float64
4   MXY         20 non-null    float64
5   NMHC        20 non-null    float64
6   NO_2        20 non-null    float64
7   NOx         20 non-null    float64
8   OXY         20 non-null    float64
9   O_3         20 non-null    float64
10  PM10        20 non-null    float64
11  PM25        20 non-null    float64
12  PXY         20 non-null    float64
13  SO_2        20 non-null    float64
14  TCH         20 non-null    float64
15  TOL         20 non-null    float64
16  station     20 non-null    int64
dtypes: float64(15), int64(1), object(1)
memory usage: 2.8+ KB
```

```
In [570]: df.describe()
```

```
Out[570]:
```

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	
count	20.000000	20.000000	20.000000	20.000000	20.000000	20.000000	20.000000	20.000000	2
mean	1.210500	0.38700	1.635000	2.736500	0.08800	60.458500	80.153500	1.611500	5
std	0.885357	0.18818	1.405394	2.382576	0.06161	30.317645	53.404650	1.152641	1
min	0.260000	0.14000	0.270000	0.390000	0.02000	24.870001	25.530001	0.380000	2
25%	0.290000	0.27000	0.452500	0.722500	0.05000	36.297500	37.044999	0.732500	3
50%	1.150000	0.34500	1.390000	2.390000	0.06000	49.594999	58.955000	1.420000	5
75%	1.680000	0.47750	2.345000	3.975000	0.15000	78.012499	112.849998	2.072500	6
max	3.240000	0.74000	5.550000	9.720000	0.21000	115.800003	189.899994	5.040000	8




```
In [571]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```



```
In [572]: sns.pairplot(df1[0:50])
```

```
Out[572]: <seaborn.axisgrid.PairGrid at 0x1918e072eb0>
```




```
In [573]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

```
Out[573]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [574]: sns.heatmap(df1.corr())
```

```
Out[574]: <AxesSubplot:>
```



```
In [575]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [576]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [577]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[577]: LinearRegression()

```
In [578]: lr.intercept_
```

Out[578]: 28078611.61393362


```
In [579]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[579]:

	Co-efficient
BEN	46.289259
CO	-222.315735
EBE	42.375662
MXY	-84.578916
NMHC	799.168387
NO_2	3.605769
NOx	-1.045082
OXY	58.701244
O_3	2.973096
PM10	-0.831940
PXY	0.828973
SO_2	13.869036
TCH	44.506946
TOL	3.758407

```
In [580]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[580]: <matplotlib.collections.PathCollection at 0x19198e7bca0>



```
In [581]: lr.score(x_test,y_test)
```

Out[581]: -8.913696645528033

```
In [582]: lr.score(x_train,y_train)
```

Out[582]: 1.0

```
In [583]: from sklearn.linear_model import Ridge,Lasso
```

```
In [584]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[584]: Ridge(alpha=10)

```
In [585]: rr.score(x_test,y_test)
```

Out[585]: -25.680273390727546

```
In [586]: rr.score(x_train,y_train)
```

Out[586]: 0.4757690329552654

```
In [587]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[587]: Lasso(alpha=10)

```
In [588]: la.score(x_train,y_train)
```

```
Out[588]: 0.397166976827539
```

```
In [589]: la.score(x_test,y_test)
```

```
Out[589]: -23.39084298600196
```

```
In [590]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[590]: ElasticNet()
```

```
In [591]: en.coef_
```

```
Out[591]: array([ 0.35276679,  0.59873896, -0.7033591 ,  5.17705657,  2.24177463,  
                -2.27215567, -0.13267477,  3.76843052, -1.97925011,  0.99602759,  
                3.50473413, -3.50906271,  0.42418908,  5.22961332])
```

```
In [592]: en.intercept_
```

```
Out[592]: 28079231.929840643
```

```
In [593]: prediction=en.predict(x_test)
```

```
In [594]: en.score(x_test,y_test)
```

```
Out[594]: -25.999562917054273
```

```
In [595]: from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
39.17938381433487  
2186.964596281396  
46.7649932778932
```

```
In [596]: from sklearn.linear_model import LogisticRegression
```

```
In [597]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O  
PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```

```
In [598]: feature_matrix.shape
```

```
Out[598]: (20, 14)
```

```
In [599]: target_vector.shape
```

```
Out[599]: (20,)
```

```
In [600]: from sklearn.preprocessing import StandardScaler
```

```
In [601]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [602]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[602]: LogisticRegression(max_iter=10000)
```

```
In [603]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [604]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079099]
```

```
In [605]: logr.score(fs,target_vector)
```

```
Out[605]: 1.0
```

```
In [606]: logr.predict_proba(observation)[0][0]
```

```
Out[606]: 0.23544436221308845
```

```
In [607]: logr.predict_proba(observation)
```

```
Out[607]: array([[2.35444362e-01, 1.79916389e-11, 7.64555638e-01]])
```

```
In [608]: from sklearn.ensemble import RandomForestClassifier
```

```
In [609]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[609]: RandomForestClassifier()
```

```
In [610]: parameters={'max_depth':[1,2,3,4,5],
  'min_samples_leaf':[5,10,15,20,25],
  'n_estimators':[10,20,30,40,50]}
```

```
In [611]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
grid_search.fit(x_train,y_train)
```

```
Out[611]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
  param_grid={'max_depth': [1, 2, 3, 4, 5],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
  scoring='accuracy')
```


```
In [612]: grid_search.best_score_
```

```
Out[612]: 0.42857142857142855
```

```
In [613]: rfc_best=grid_search.best_estimator_
```

```
In [614]: from sklearn.tree import plot_tree
plt.figure(figsize=(50,5))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'],
```

```
Out[614]: [Text(1395.0, 203.85000000000002, 'EBE <= 1.39\ngini = 0.663\nsamples = 11\nv
  alue = [4, 5, 5]\nclass = b'),
  Text(697.5, 67.94999999999999, 'gini = 0.408\nsamples = 5\nvalue = [0, 5, 2]
  \nclass = b'),
  Text(2092.5, 67.94999999999999, 'gini = 0.49\nsamples = 6\nvalue = [4, 0, 3]
  \nclass = a')]
```



Conclusion

Linear Regression =1.0

Ridge Regression =0.8505902210559991

Lasso Regression =0.7818422293632583

ElasticNet Regression =0.6639225233261004

Logistic Regression =0.8991524766814838

Randomforest =0.3571428571428571

Logistic Regression is suitable for this dataset

In []: