

mk 3/08/2023

```
In [147]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

In [148]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2010.csv")
df
```

Out[148]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	I
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	I
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	I
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19.410
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24.670
...	...	...	...	...	...	...	...	...	...	...	...
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999	I
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN	51.259
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000	I
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999	I
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998	47.150

209448 rows × 12 columns

```
In [149]: df=df.dropna()
```

```
In [150]: df=df.head(100)
```

```
In [151]: df.columns
```

```
Out[151]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
               dtype='object')
```

```
In [152]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 100 entries, 11 to 1199  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   date        100 non-null    object  
1   BEN         100 non-null    float64  
2   CO          100 non-null    float64  
3   EBE         100 non-null    float64  
4   MXY         100 non-null    float64  
5   NMHC        100 non-null    float64  
6   NO_2        100 non-null    float64  
7   NOx         100 non-null    float64  
8   OXY         100 non-null    float64  
9   O_3         100 non-null    float64  
10  PM10        100 non-null    float64  
11  PM25        100 non-null    float64  
12  PXY         100 non-null    float64  
13  SO_2        100 non-null    float64  
14  TCH         100 non-null    float64  
15  TOL         100 non-null    float64  
16  station     100 non-null    int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 14.1+ KB
```

```
In [153]: data=df[['BEN', 'TOL', 'PXY']]
data
```


Out[153]:

	BEN	TOL	PXY
11	0.78	1.99	0.81
23	0.70	2.62	0.81
35	0.58	0.84	0.81
47	0.33	1.21	0.81
59	0.38	0.49	0.79
...	...	...	...
1151	0.44	1.66	0.53
1163	0.45	0.86	0.57
1175	0.40	1.48	0.57
1187	0.45	0.91	0.56
1199	0.42	1.60	0.56

100 rows × 3 columns


```
In [154]: data.plot.line(subplots=True)
```

Out[154]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)



```
In [155]: data.plot.line()
```


```
Out[155]: <AxesSubplot:>
```



```
In [156]: b=data[0:50]
```


```
In [157]: b.plot.bar()
```

```
Out[157]: <AxesSubplot:>
```




```
In [158]: data.plot.hist()
```

```
Out[158]: <AxesSubplot:ylabel='Frequency'>
```




```
In [159]: data.plot.area()
```

```
Out[159]: <AxesSubplot:>
```




```
In [160]: data.plot.box()
```

```
Out[160]: <AxesSubplot:>
```




```
In [161]: b.plot.pie(y='BEN' )
```

```
Out[161]: <AxesSubplot:ylabel='BEN'>
```



```
In [162]: data.plot.scatter(x='TOL', y='BEN')
```

```
Out[162]: <AxesSubplot:xlabel='TOL', ylabel='BEN'>
```



```
In [163]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100 entries, 11 to 1199
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        100 non-null   object
 1   BEN         100 non-null   float64
 2   CO          100 non-null   float64
 3   EBE         100 non-null   float64
 4   MXY         100 non-null   float64
 5   NMHC        100 non-null   float64
 6   NO_2        100 non-null   float64
 7   NOx         100 non-null   float64
 8   OXY         100 non-null   float64
 9   O_3         100 non-null   float64
10  PM10        100 non-null   float64
11  PM25        100 non-null   float64
12  PXY         100 non-null   float64
13  SO_2        100 non-null   float64
14  TCH         100 non-null   float64
15  TOL         100 non-null   float64
16  station     100 non-null   int64
dtypes: float64(15), int64(1), object(1)
memory usage: 14.1+ KB
```



```
In [164]: df.describe()
```

Out[164]:

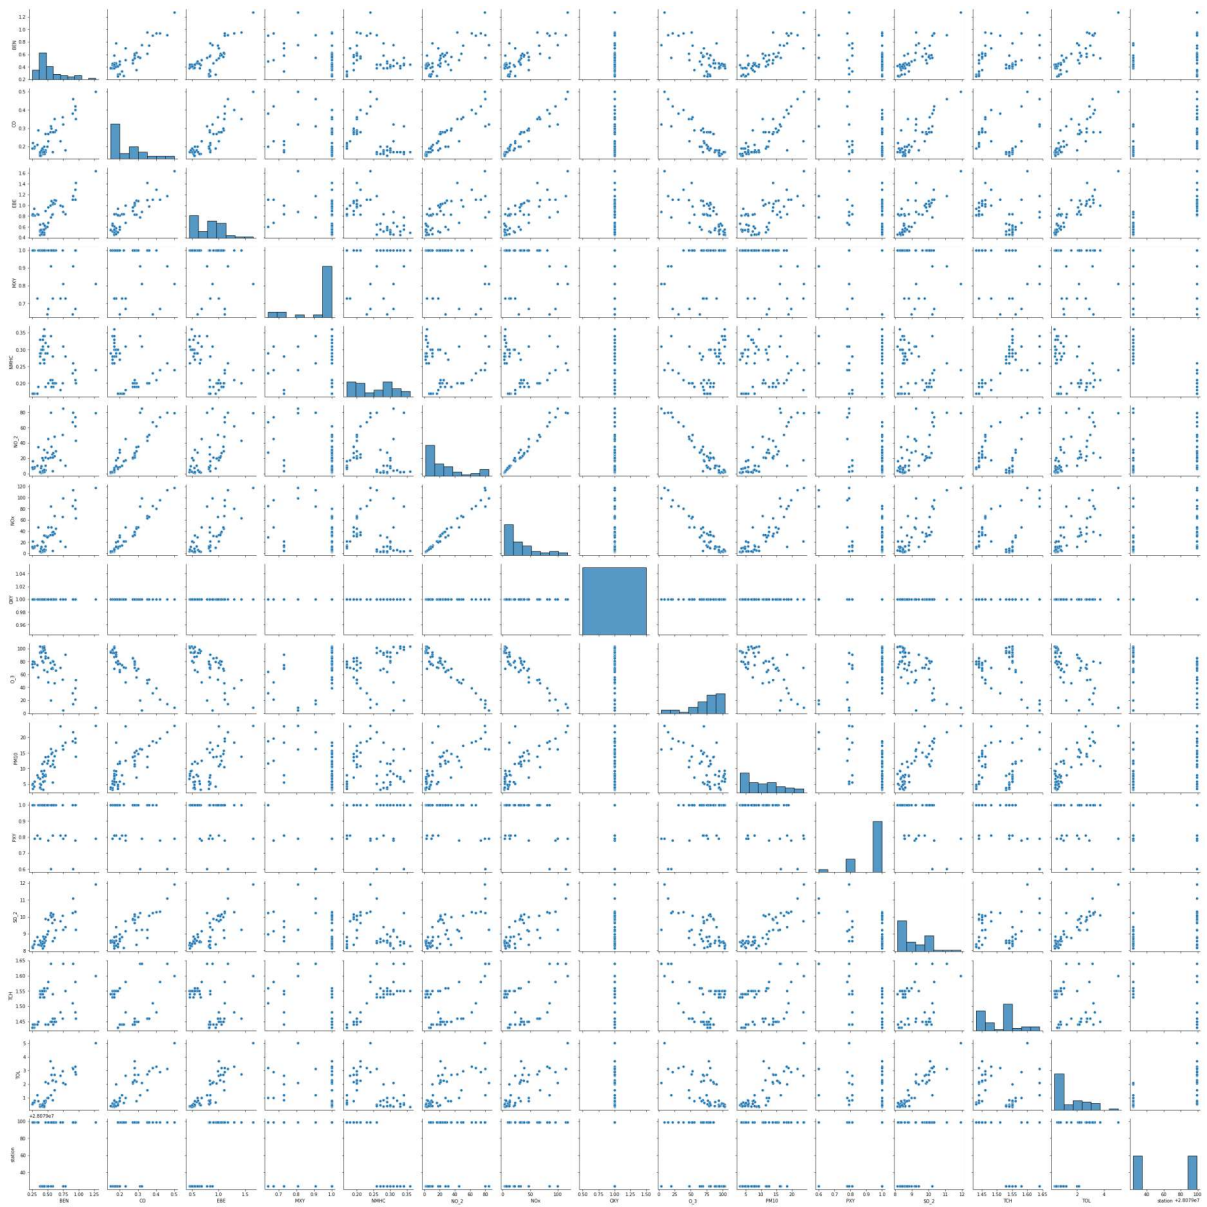
	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	
<b>count</b>	100.000000	100.000000	100.000000	100.000000	100.000000	100.000000	100.000000	100.0
<b>mean</b>	0.597600	0.258900	0.920100	0.842200	0.275100	29.911300	38.454500	1.0
<b>std</b>	0.254626	0.083217	0.352931	0.284575	0.079181	20.908565	27.822058	0.2
<b>min</b>	0.250000	0.150000	0.340000	0.380000	0.170000	1.290000	2.780000	0.5
<b>25%</b>	0.440000	0.190000	0.655000	0.610000	0.210000	16.910000	20.057500	1.0
<b>50%</b>	0.555000	0.240000	0.880000	0.935000	0.260000	27.100000	33.234999	1.0
<b>75%</b>	0.660000	0.310000	1.102500	1.000000	0.330000	40.415001	50.852500	1.0
<b>max</b>	1.630000	0.500000	2.180000	1.780000	0.470000	84.629997	117.300003	1.8



```
In [165]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [166]: sns.pairplot(df1[0:50])
```

```
Out[166]: <seaborn.axisgrid.PairGrid at 0x1d8f8c34760>
```




```
In [167]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

```
Out[167]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [168]: sns.heatmap(df1.corr())
```

```
Out[168]: <AxesSubplot:>
```



```
In [169]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [170]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [171]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[171]: LinearRegression()

```
In [172]: lr.intercept_
```

Out[172]: 28079197.48878474


```
In [173]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[173]:

	Co-efficient
BEN	-65.477903
CO	562.027959
EBE	27.200743
MXY	-10.248607
NMHC	-111.487476
NO_2	-0.741432
NOx	-0.447744
OXY	16.357217
O_3	-0.235314
PM10	0.116936
PXY	-3.133652
SO_2	-3.087070
TCH	-106.991085
TOL	2.717068

```
In [174]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[174]: <matplotlib.collections.PathCollection at 0x1d886e84640>



```
In [175]: lr.score(x_test,y_test)
```

Out[175]: 0.8896517498224122

```
In [176]: lr.score(x_train,y_train)
```

Out[176]: 0.949306857592614

```
In [177]: from sklearn.linear_model import Ridge,Lasso
```

```
In [178]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[178]: Ridge(alpha=10)

```
In [179]: rr.score(x_test,y_test)
```

Out[179]: 0.4697103996800025

```
In [180]: rr.score(x_train,y_train)
```

Out[180]: 0.6520503785699647

```
In [181]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[181]: Lasso(alpha=10)

```
In [182]: la.score(x_train,y_train)
```

```
Out[182]: 0.4692794984392259
```

```
In [183]: la.score(x_test,y_test)
```

```
Out[183]: 0.36782468523872713
```

```
In [184]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[184]: ElasticNet()
```

```
In [185]: en.coef_
```

```
Out[185]: array([-1.92676588,  0.          ,  1.79373525,  3.68565428, -2.03130748,  
                -3.14780078,  2.35863109, -0.          , -0.35659703, -1.21336016,  
                -0.39958075,  2.97652582, -1.38750317,  6.95817564])
```

```
In [186]: en.intercept_
```

```
Out[186]: 28079057.24657358
```

```
In [187]: prediction=en.predict(x_test)
```

```
In [188]: en.score(x_test,y_test)
```

```
Out[188]: 0.3895631825806143
```

```
In [189]: from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
25.601336059346796
```

```
843.1658540605265
```

```
29.037318300086294
```

```
In [190]: from sklearn.linear_model import LogisticRegression
```

```
In [191]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O  
PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```

```
In [192]: feature_matrix.shape
```

```
Out[192]: (100, 14)
```

```
In [193]: target_vector.shape
```

```
Out[193]: (100,)
```

```
In [194]: from sklearn.preprocessing import StandardScaler
```

```
In [195]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [196]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[196]: LogisticRegression(max_iter=10000)
```

```
In [197]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [198]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079024]
```

```
In [199]: logr.score(fs,target_vector)
```

```
Out[199]: 1.0
```

```
In [200]: logr.predict_proba(observation)[0][0]
```

```
Out[200]: 0.9993168651234103
```

```
In [201]: logr.predict_proba(observation)
```

```
Out[201]: array([[9.99316865e-01, 6.83134877e-04]])
```

```
In [202]: from sklearn.ensemble import RandomForestClassifier
```

```
In [203]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[203]: RandomForestClassifier()
```

```
In [204]: parameters={'max_depth':[1,2,3,4,5],
  'min_samples_leaf':[5,10,15,20,25],
  'n_estimators':[10,20,30,40,50]}
```

```
In [205]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
grid_search.fit(x_train,y_train)
```

```
Out[205]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
  param_grid={'max_depth': [1, 2, 3, 4, 5],
  'min_samples_leaf': [5, 10, 15, 20, 25],
  'n_estimators': [10, 20, 30, 40, 50]},
  scoring='accuracy')
```


```
In [206]: grid_search.best_score_
```

```
Out[206]: 1.0
```

```
In [207]: rfc_best=grid_search.best_estimator_
```

```
In [208]: from sklearn.tree import plot_tree
plt.figure(figsize=(50,5))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'],
```

```
Out[208]: [Text(1395.0, 203.85000000000002, 'TCH <= 1.525\ngini = 0.474\nsamples = 44\n
value = [43, 27]\nclass = a'),
  Text(697.5, 67.94999999999999, 'gini = 0.0\nsamples = 14\nvalue = [0, 21]\nc
lass = b'),
  Text(2092.5, 67.94999999999999, 'gini = 0.215\nsamples = 30\nvalue = [43, 6]
\nclass = a')]
```



Conclusion

Linear Regression =0.949306857592614

Ridge Regression =0.6520503785699647

Lasso Regression =0.4692794984392259

ElasticNet Regression =0.3895631825806143

Logistic Regression =0.9993168651234103

Randomforest =1.0

Randomforest is suitable for this dataset



In [ ]: