

mk 3/08/2023

```
In [316]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [317]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2011-11-01-01:00:00.csv")
df
```

Out[317]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	NaN	2
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7	2
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	7.2	2
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	NaN	2
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	NaN	2
...
209923	2011-09-01 00:00:00	NaN	0.2	NaN	NaN	5.0	19.0	44.0	NaN	NaN	NaN	NaN	NaN	2
209924	2011-09-01 00:00:00	NaN	0.1	NaN	NaN	6.0	29.0	NaN	11.0	NaN	7.0	NaN	NaN	2
209925	2011-09-01 00:00:00	NaN	NaN	NaN	0.23	1.0	21.0	28.0	NaN	NaN	NaN	1.44	NaN	2
209926	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	15.0	48.0	NaN	NaN	NaN	NaN	NaN	2
209927	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	4.0	33.0	38.0	13.0	NaN	NaN	NaN	NaN	2

209928 rows × 14 columns



```
In [318]: df=df.dropna()
```

```
In [319]: df=df.head(800)
```

```
In [320]: df.columns
```

```
Out[320]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
                'SO_2', 'TCH', 'TOL', 'station'],  
               dtype='object')
```

```
In [321]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 800 entries, 1 to 9721  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   date        800 non-null    object  
1   BEN         800 non-null    float64  
2   CO          800 non-null    float64  
3   EBE         800 non-null    float64  
4   NMHC        800 non-null    float64  
5   NO          800 non-null    float64  
6   NO_2        800 non-null    float64  
7   O_3         800 non-null    float64  
8   PM10        800 non-null    float64  
9   PM25        800 non-null    float64  
10  SO_2        800 non-null    float64  
11  TCH         800 non-null    float64  
12  TOL         800 non-null    float64  
13  station     800 non-null    int64  
dtypes: float64(12), int64(1), object(1)  
memory usage: 93.8+ KB
```

```
In [322]: data=df[['BEN', 'TOL']]
data
```


Out[322]:

	BEN	TOL
1	2.5	8.7
6	0.7	1.7
25	1.8	7.4
30	1.0	2.9
49	1.3	6.2
...
9673	2.4	8.1
9678	1.0	3.4
9697	2.8	9.1
9702	1.2	3.8
9721	1.9	6.2

800 rows × 2 columns


```
In [323]: data.plot.line(subplots=True)
```

Out[323]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



```
In [324]: data.plot.line()
```


```
Out[324]: <AxesSubplot:>
```



```
In [325]: b=data[0:50]
```


```
In [326]: b.plot.bar()
```

```
Out[326]: <AxesSubplot:>
```




```
In [327]: data.plot.hist()
```

```
Out[327]: <AxesSubplot:ylabel='Frequency'>
```



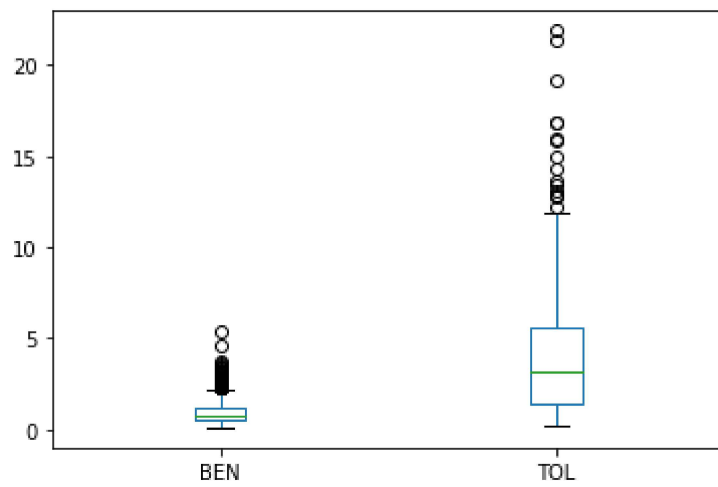
```
In [328]: data.plot.area()
```

```
Out[328]: <AxesSubplot:>
```




```
In [329]: data.plot.box()
```

```
Out[329]: <AxesSubplot:>
```




```
In [330]: b.plot.pie(y='BEN' )
```

```
Out[330]: <AxesSubplot:ylabel='BEN'>
```



```
In [331]: data.plot.scatter(x='TOL', y='BEN')
```

```
Out[331]: <AxesSubplot:xlabel='TOL', ylabel='BEN'>
```



```
In [332]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 800 entries, 1 to 9721
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        800 non-null    object
 1   BEN         800 non-null    float64
 2   CO          800 non-null    float64
 3   EBE         800 non-null    float64
 4   NMHC        800 non-null    float64
 5   NO          800 non-null    float64
 6   NO_2        800 non-null    float64
 7   O_3         800 non-null    float64
 8   PM10        800 non-null    float64
 9   PM25        800 non-null    float64
10   SO_2        800 non-null    float64
11   TCH         800 non-null    float64
12   TOL         800 non-null    float64
13   station     800 non-null    int64
dtypes: float64(12), int64(1), object(1)
memory usage: 93.8+ KB
```



```
In [333]: df.describe()
```

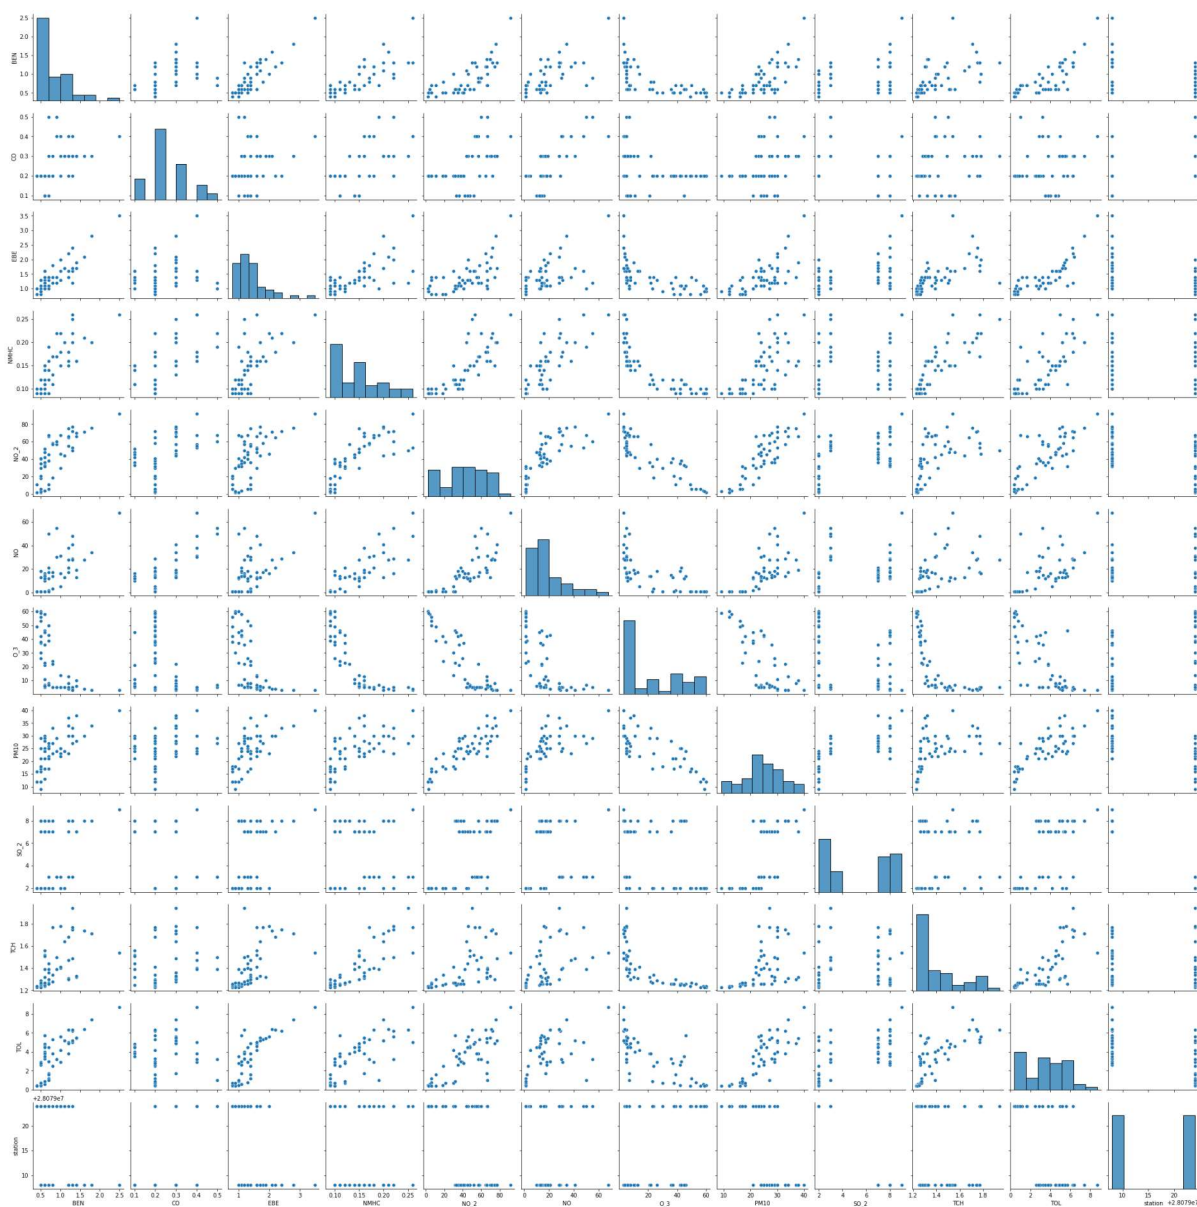
Out[333]:

	BEN	CO	EBE	NMHC	NO	NO_2	O_3	
count	800.00000	800.000000	800.000000	800.000000	800.000000	800.000000	800.000000	800.000000
mean	0.94225	0.296750	1.475750	0.161412	25.791250	42.945000	22.772500	18.760000
std	0.68275	0.162704	1.041041	0.053453	35.160995	25.994213	17.267917	10.430000
min	0.10000	0.100000	0.300000	0.080000	1.000000	1.000000	2.000000	2.000000
25%	0.50000	0.200000	0.800000	0.120000	2.000000	22.000000	6.000000	11.000000
50%	0.70000	0.300000	1.100000	0.150000	14.000000	43.500000	19.000000	18.000000
75%	1.20000	0.400000	1.700000	0.190000	36.000000	61.000000	37.000000	25.000000
max	5.40000	1.900000	8.200000	0.490000	365.000000	175.000000	64.000000	89.000000

```
In [334]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3', 'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [335]: sns.pairplot(df1[0:50])
```

```
Out[335]: <seaborn.axisgrid.PairGrid at 0x1d88a3f7430>
```




```
In [336]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

```
Out[336]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [337]: sns.heatmap(df1.corr())
```

```
Out[337]: <AxesSubplot:>
```



```
In [338]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',  
                'PM10', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
In [339]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [340]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[340]: LinearRegression()

```
In [341]: lr.intercept_
```

Out[341]: 28079021.82965764


```
In [342]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[342]:

	Co-efficient
BEN	-0.034243
CO	22.339344
EBE	-0.741200
NMHC	8.754248
NO_2	-0.075296
NO	-0.012261
O_3	0.032894
PM10	0.051771
SO_2	-2.097788
TCH	0.133014
TOL	0.149400

```
In [343]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[343]: <matplotlib.collections.PathCollection at 0x1d893640550>



```
In [344]: lr.score(x_test,y_test)
```

Out[344]: 0.8895515441584051

```
In [345]: lr.score(x_train,y_train)
```

Out[345]: 0.9063734845521797

```
In [346]: from sklearn.linear_model import Ridge,Lasso
```

```
In [347]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[347]: Ridge(alpha=10)

```
In [348]: rr.score(x_test,y_test)
```

Out[348]: 0.8687242338146928

```
In [349]: rr.score(x_train,y_train)
```

Out[349]: 0.8888695976199741

```
In [350]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[350]: Lasso(alpha=10)

```
In [351]: la.score(x_train,y_train)
```

```
Out[351]: 0.6073128952517146
```

```
In [352]: la.score(x_test,y_test)
```

```
Out[352]: 0.5343332803911878
```

```
In [353]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[353]: ElasticNet()
```

```
In [354]: en.coef_
```

```
Out[354]: array([ 0.          ,  0.          ,  0.          ,  0.          , -0.06491198,
                0.06997964, -0.          ,  0.03644468, -2.11361822,  0.          ,
                0.          ])
```

```
In [355]: en.intercept_
```

```
Out[355]: 28079028.09886945
```

```
In [356]: prediction=en.predict(x_test)
```

```
In [357]: en.score(x_test,y_test)
```

```
Out[357]: 0.8407158041429669
```

```
In [358]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
2.5001350118933865
```

```
10.136846224341589
```

```
3.183841425753109
```

```
In [359]: from sklearn.linear_model import LogisticRegression
```

```
In [360]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',
                             'PM10', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [361]: feature_matrix.shape
```

```
Out[361]: (800, 11)
```

```
In [362]: target_vector.shape
```

```
Out[362]: (800,)
```

```
In [363]: from sklearn.preprocessing import StandardScaler
```

```
In [364]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [365]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[365]: LogisticRegression(max_iter=10000)
```

```
In [370]: observation=[[1,2,3,4,5,6,7,8,9,10,11]]
```

```
In [371]: prediction=logr.predict(observation)
print(prediction)
```

```
[28079008]
```

```
In [372]: logr.score(fs,target_vector)
```

```
Out[372]: 1.0
```

```
In [373]: logr.predict_proba(observation)[0][0]
```

```
Out[373]: 1.0
```

```
In [374]: logr.predict_proba(observation)
```

```
Out[374]: array([[1.00000000e+00, 4.27828389e-22]])
```

```
In [375]: from sklearn.ensemble import RandomForestClassifier
```

```
In [376]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[376]: RandomForestClassifier()
```

```
In [377]: parameters={'max_depth':[1,2,3,4,5],
  'min_samples_leaf':[5,10,15,20,25],
  'n_estimators':[10,20,30,40,50]}
```

```
In [378]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
grid_search.fit(x_train,y_train)
```

```
Out[378]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
  param_grid={'max_depth': [1, 2, 3, 4, 5],
  'min_samples_leaf': [5, 10, 15, 20, 25],
  'n_estimators': [10, 20, 30, 40, 50]},
  scoring='accuracy')
```


```
In [379]: grid_search.best_score_
```

```
Out[379]: 0.9964285714285714
```

```
In [380]: rfc_best=grid_search.best_estimator_
```

```
In [381]: from sklearn.tree import plot_tree
plt.figure(figsize=(50,5))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'],
```

```
Out[381]: [Text(1395.0, 203.85000000000002, 'EBE <= 1.45\ngini = 0.5\nsamples = 352\nva
  lue = [278, 282]\nclass = b'),
  Text(697.5, 67.94999999999999, 'gini = 0.448\nsamples = 235\nvalue = [132, 2
  58]\nclass = b'),
  Text(2092.5, 67.94999999999999, 'gini = 0.242\nsamples = 117\nvalue = [146,
  24]\nclass = a')]
```



Conclusion

Linear Regression =0.9063734845521797

Ridge Regression =0.8888695976199741

Lasso Regression =0.6073128952517146

ElasticNet Regression =0.8407158041429669

Logistic Regression =1.0

Randomforest =0.9964285714285714

Logistic Regression is suitable for this dataset

In []: