

mk 3/08/2023

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.preprocessing import StandardScaler
7 import re
8 from sklearn.datasets import load_digits
9 from sklearn.model_selection import train_test_split
```

In [4]:

```
1 a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_yea
2 a
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2013-11-01 01:00:00	NaN	0.6	NaN	NaN	135.0	74.0	NaN	NaN	NaN	7.0	NaN	NaN	2
1	2013-11-01 01:00:00	1.5	0.5	1.3	NaN	71.0	83.0	2.0	23.0	16.0	12.0	NaN	8.3	2
2	2013-11-01 01:00:00	3.9	NaN	2.8	NaN	49.0	70.0	NaN	NaN	NaN	NaN	NaN	9.0	2
3	2013-11-01 01:00:00	NaN	0.5	NaN	NaN	82.0	87.0	3.0	NaN	NaN	NaN	NaN	NaN	2
4	2013-11-01 01:00:00	NaN	NaN	NaN	NaN	242.0	111.0	2.0	NaN	NaN	12.0	NaN	NaN	2
...	
209875	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	8.0	39.0	52.0	NaN	NaN	NaN	NaN	NaN	2
209876	2013-03-01 00:00:00	NaN	0.4	NaN	NaN	1.0	11.0	NaN	6.0	NaN	2.0	NaN	NaN	2
209877	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	4.0	75.0	NaN	NaN	NaN	NaN	NaN	2
209878	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	11.0	52.0	NaN	NaN	NaN	NaN	NaN	2
209879	2013-03-01 00:00:00	NaN	NaN	NaN	NaN	1.0	10.0	75.0	3.0	NaN	NaN	NaN	NaN	2
209880 rows × 14 columns														

In [5]: 1 a.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209880 entries, 0 to 209879
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209880 non-null object
1   BEN         50462 non-null  float64
2   CO          87018 non-null  float64
3   EBE         50463 non-null  float64
4   NMHC        25935 non-null  float64
5   NO          209108 non-null float64
6   NO_2        209108 non-null float64
7   O_3         121858 non-null float64
8   PM10        104339 non-null float64
9   PM25        51980 non-null  float64
10  SO_2        86970 non-null  float64
11  TCH         25935 non-null  float64
12  TOL         50317 non-null  float64
13  station     209880 non-null int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

In [6]:

1

b=a.fillna(value=104)

2

b

Out[6]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	T
0	2013-11-01 01:00:00	104.0	0.6	104.0	104.0	135.0	74.0	104.0	104.0	104.0	7.0	104.0	10
1	2013-11-01 01:00:00	1.5	0.5	1.3	104.0	71.0	83.0	2.0	23.0	16.0	12.0	104.0	
2	2013-11-01 01:00:00	3.9	104.0	2.8	104.0	49.0	70.0	104.0	104.0	104.0	104.0	104.0	
3	2013-11-01 01:00:00	104.0	0.5	104.0	104.0	82.0	87.0	3.0	104.0	104.0	104.0	104.0	10
4	2013-11-01 01:00:00	104.0	104.0	104.0	104.0	242.0	111.0	2.0	104.0	104.0	12.0	104.0	10
...	
209875	2013-03-01 00:00:00	104.0	0.4	104.0	104.0	8.0	39.0	52.0	104.0	104.0	104.0	104.0	10
209876	2013-03-01 00:00:00	104.0	0.4	104.0	104.0	1.0	11.0	104.0	6.0	104.0	2.0	104.0	10
209877	2013-03-01 00:00:00	104.0	104.0	104.0	104.0	2.0	4.0	75.0	104.0	104.0	104.0	104.0	10
209878	2013-03-01 00:00:00	104.0	104.0	104.0	104.0	2.0	11.0	52.0	104.0	104.0	104.0	104.0	10
209879	2013-03-01 00:00:00	104.0	104.0	104.0	104.0	1.0	10.0	75.0	3.0	104.0	104.0	104.0	10
209880 rows × 14 columns													

In [7]:

1

b.columns

Out[7]:

Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
'SO_2', 'TCH', 'TOL', 'station'],
 dtype='object')

In [8]:

1

c=b.head(10000)

2

c

Out[8]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOI
0	2013-11-01 01:00:00	104.0	0.6	104.0	104.0	135.0	74.0	104.0	104.0	104.0	7.0	104.0	104.0
1	2013-11-01 01:00:00	1.5	0.5	1.3	104.0	71.0	83.0	2.0	23.0	16.0	12.0	104.0	8.0
2	2013-11-01 01:00:00	3.9	104.0	2.8	104.0	49.0	70.0	104.0	104.0	104.0	104.0	104.0	9.0
3	2013-11-01 01:00:00	104.0	0.5	104.0	104.0	82.0	87.0	3.0	104.0	104.0	104.0	104.0	104.0
4	2013-11-01 01:00:00	104.0	104.0	104.0	104.0	242.0	111.0	2.0	104.0	104.0	12.0	104.0	104.0
...
9995	2013-11-18 09:00:00	104.0	0.7	104.0	104.0	93.0	57.0	4.0	104.0	104.0	104.0	104.0	104.0
9996	2013-11-18 09:00:00	104.0	104.0	104.0	104.0	138.0	69.0	104.0	23.0	104.0	6.0	104.0	104.0
9997	2013-11-18 09:00:00	104.0	104.0	104.0	104.0	168.0	64.0	104.0	22.0	15.0	104.0	104.0	104.0
9998	2013-11-18 09:00:00	104.0	104.0	104.0	104.0	110.0	89.0	104.0	22.0	16.0	104.0	104.0	104.0
9999	2013-11-18 09:00:00	104.0	104.0	104.0	104.0	53.0	42.0	2.0	104.0	104.0	104.0	104.0	104.0

10000 rows × 14 columns

```
In [9]: 1 d=c[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
2         'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
3 d
```


Out[9]:

	BEN	CO	EBE	NMHC	NO_2	O_3	PM10	SO_2	TCH	TOL	station
0	104.0	0.6	104.0	104.0	74.0	104.0	104.0	7.0	104.0	104.0	28079004
1	1.5	0.5	1.3	104.0	83.0	2.0	23.0	12.0	104.0	8.3	28079008
2	3.9	104.0	2.8	104.0	70.0	104.0	104.0	104.0	104.0	9.0	28079011
3	104.0	0.5	104.0	104.0	87.0	3.0	104.0	104.0	104.0	104.0	28079016
4	104.0	104.0	104.0	104.0	111.0	2.0	104.0	12.0	104.0	104.0	28079017
...
9995	104.0	0.7	104.0	104.0	57.0	4.0	104.0	104.0	104.0	104.0	28079039
9996	104.0	104.0	104.0	104.0	69.0	104.0	23.0	6.0	104.0	104.0	28079040
9997	104.0	104.0	104.0	104.0	64.0	104.0	22.0	104.0	104.0	104.0	28079047
9998	104.0	104.0	104.0	104.0	89.0	104.0	22.0	104.0	104.0	104.0	28079048
9999	104.0	104.0	104.0	104.0	42.0	2.0	104.0	104.0	104.0	104.0	28079049

10000 rows × 11 columns

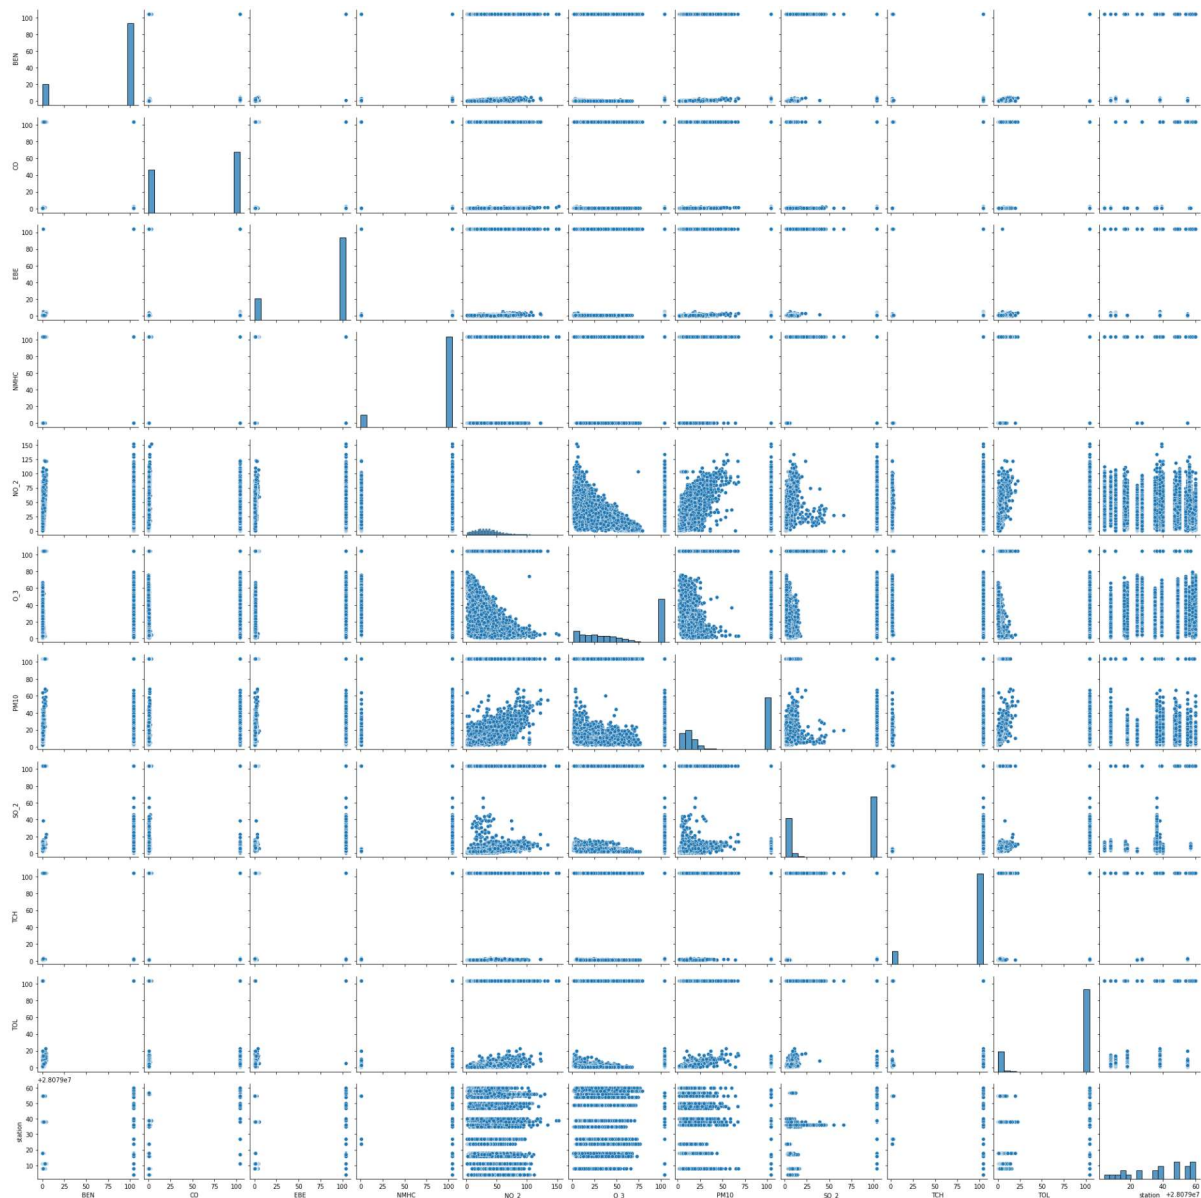
```
In [10]: 1 d.plot.line()
```

Out[10]: <AxesSubplot:>



```
In [11]: 1 sns.pairplot(d)
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x1f70ad5bd90>
```



```
In [12]: 1 x=d[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2']]
2 y=d['TCH']
```

```
In [13]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: 1 from sklearn.linear_model import LinearRegression
2 lr=LinearRegression()
3
```

```
Out[14]: LinearRegression()
```

```
In [47]: 1 lr.fit(x_train,y_train)
```

```
Out[47]: LinearRegression()
```

```
In [15]: 1 print(lr.intercept_)
```

```
1.2479844840821386
```


```
In [16]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
2 coeff
```

```
Out[16]:
```

Co-efficient	
BEN	-0.000969
CO	0.000221
EBE	0.000703
NMHC	0.987915
NO_2	0.000508

```
In [17]: 1 prediction=lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x1f714444850>
```



```
In [18]: 1 print(lr.score(x_test,y_test))
```

```
0.9999953135331272
```

```
In [48]: 1 print(lr.score(x_train,y_train))
```

```
0.9999954606932426
```

```
In [19]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

Out[20]: Ridge(alpha=10)

```
In [21]: 1 rr.score(x_test,y_test)
```

Out[21]: 0.9999953138694069

```
In [22]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

Out[22]: Lasso(alpha=10)

```
In [23]: 1 la.score(x_test,y_test)
```

Out[23]: 0.9999189899159493


```
In [25]: 1 a1=b.head(7000)
          2 a1
```

Out[25]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOI
0	2013-11-01 01:00:00	104.0	0.6	104.0	104.0	135.0	74.0	104.0	104.0	104.0	7.0	104.0	104.0
1	2013-11-01 01:00:00	1.5	0.5	1.3	104.0	71.0	83.0	2.0	23.0	16.0	12.0	104.0	8.0
2	2013-11-01 01:00:00	3.9	104.0	2.8	104.0	49.0	70.0	104.0	104.0	104.0	104.0	104.0	9.0
3	2013-11-01 01:00:00	104.0	0.5	104.0	104.0	82.0	87.0	3.0	104.0	104.0	104.0	104.0	104.0
4	2013-11-01 01:00:00	104.0	104.0	104.0	104.0	242.0	111.0	2.0	104.0	104.0	12.0	104.0	104.0
...
6995	2013-11-13 04:00:00	104.0	0.2	104.0	104.0	1.0	8.0	40.0	104.0	104.0	104.0	104.0	104.0
6996	2013-11-13 04:00:00	104.0	104.0	104.0	104.0	1.0	5.0	104.0	3.0	104.0	1.0	104.0	104.0
6997	2013-11-13 04:00:00	104.0	104.0	104.0	104.0	1.0	6.0	104.0	3.0	2.0	104.0	104.0	104.0
6998	2013-11-13 04:00:00	104.0	104.0	104.0	104.0	1.0	9.0	104.0	5.0	1.0	104.0	104.0	104.0
6999	2013-11-13 04:00:00	104.0	104.0	104.0	104.0	1.0	9.0	43.0	104.0	104.0	104.0	104.0	104.0

7000 rows × 14 columns



```
In [26]: 1 e=a1[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
          2 'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [27]: 1 f=e.iloc[:,0:14]
          2 g=e.iloc[:, -1]
```

```
In [28]: 1 h=StandardScaler().fit_transform(f)
```

```
In [29]: 1 logr=LogisticRegression(max_iter=10000)
          2 logr.fit(h,g)
```

Out[29]: LogisticRegression(max_iter=10000)

```
In [30]: 1 from sklearn.model_selection import train_test_split
          2 h_train,h_test,g_train,g_test=train_test_split(h,g,test_size=0.3)
```

```
In [31]: 1 i=[[10,20,30,40,50,60,11,22,33,44,55]]
```

```
In [32]: 1 prediction=logr.predict(i)
          2 print(prediction)
```

[28079050]

```
In [33]: 1 logr.classes_
```

Out[33]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
 28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
 28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
 28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
 dtype=int64)

```
In [34]: 1 logr.predict_proba(i)[0][0]
```

Out[34]: 0.0

```
In [35]: 1 logr.predict_proba(i)[0][1]
```

Out[35]: 0.0

```
In [36]: 1 logr.score(h_test,g_test)
```

Out[36]: 0.9576190476190476

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

Out[37]: ElasticNet()

```
In [38]: 1 print(en.coef_)
```

[-0. 0. -0. 0.98701057 0.]

```
In [39]: 1 print(en.intercept_)
```

1.3400191911879489

```
In [40]: 1 prediction=en.predict(x_test)
2 print(en.score(x_test,y_test))
```

0.9999944213935371

```
In [41]: 1 from sklearn.ensemble import RandomForestClassifier
2 rfc=RandomForestClassifier()
3 rfc.fit(h_train,g_train)
```

Out[41]: RandomForestClassifier()

```
In [42]: 1 parameters={'max_depth':[1,2,3,4,5],
2 'min_samples_leaf':[5,10,15,20,25],
3 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [43]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring=
3 grid_search.fit(h_train,g_train)
```

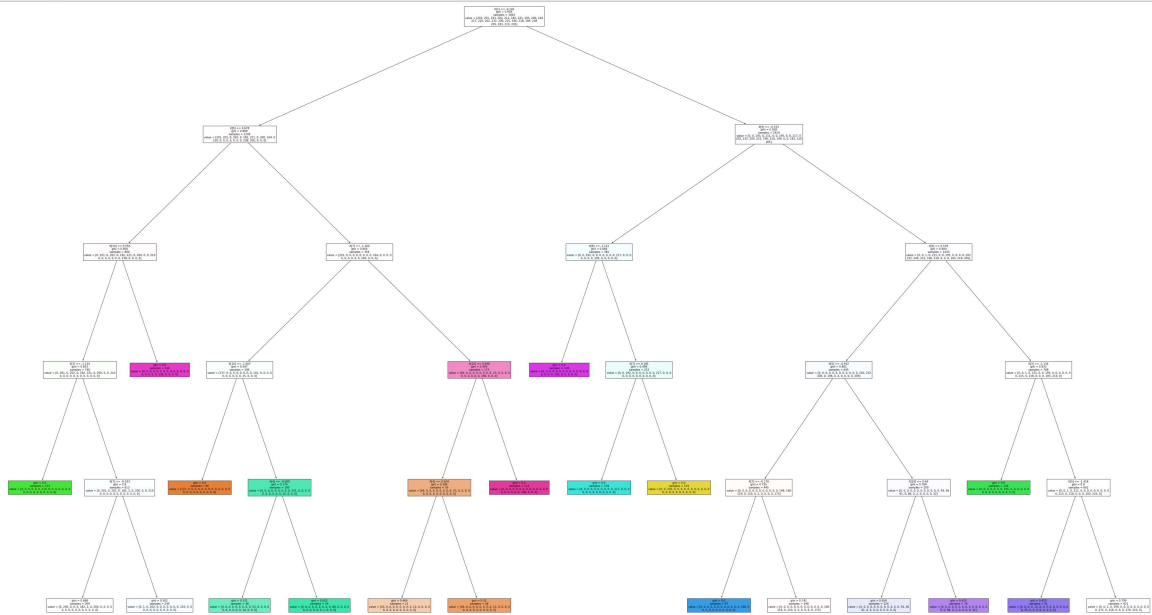
Out[43]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

```
In [44]: 1 grid_search.best_score_
```

Out[44]: 0.9985714285714286

```
In [45]: 1 rfc_best=grid_search.best_estimator_
```

```
In [46]: 1 from sklearn.tree import plot_tree
2 plt.figure(figsize=(80,50))
3 plot_tree(rfc_best.estimators_[2],filled=True)
```



Conclusion

Linear Regression=0.9999954606932426

Ridge Regression=0.9999189899159493

Lasso Regression=0.9999189899159493

ElasticNet Regression=0.9999944213935371

Logistic Regression=0.9576190476190476

Random Forest=0.9985714285714286

```
In [ ]: 1 Linear Regression is suitable in this dataset
```